

# AI Assignment #5 Report

Name: Sajjad Pakdaman Savoji

SID: 810195517

## 4. Testing Network:

*For this part the corresponding script was used.*

```
• code_sajjad python3 neural_net_tester.py simple

Training on OR data
train: 100%|
 10000/10000 [00:02<00:00, 4036.88it/s]
weights: [w1A(11.17), w2A(11.17), wA(6.93)]
Trained weights:
Weight 'w1A': 11.172709
Weight 'w2A': 11.173045
Weight 'wA': 6.933418
Testing on OR test-data
test((0.1, 0.1, 0)) returned: 0.0090735107200411206 => 0 [correct]
test((0.1, 0.9, 1)) returned: 0.9857913164734841 => 1 [correct]
test((0.9, 0.1, 1)) returned: 0.9857875715894618 => 1 [correct]
test((0.9, 0.9, 1)) returned: 0.999990107700679 => 1 [correct]
Accuracy: 1.000000
-----
Training on AND data
train: 100%|
 10000/10000 [00:02<00:00, 4036.96it/s]
weights: [w1A(10.50), w2A(10.50), wA(14.37)]
Trained weights:
Weight 'w1A': 10.099844
Weight 'w2A': 10.409561
Weight 'wA': 14.366979
Testing on AND test-data
test((0.1, 0.1, 0)) returned: 4.784254617957318e-06 => 0 [correct]
test((0.1, 0.9, 0)) returned: 0.020484406009173127 => 0 [correct]
test((0.9, 0.1, 0)) returned: 0.020484406009173127 => 0 [correct]
test((0.9, 0.9, 1)) returned: 0.9893584979736043 => 1 [correct]
Accuracy: 1.000000
```

## 5. Finite difference:

For this part member function `finite_difference(self)` was implemented for class `Network` using which correctness of derivatives implementations was tested by simple test

```
+ code_sajjad python3 neural_net_tester.py simple
-----
Training on OR data
train: 100% | 10000/10000 [00:02:00:00, 3805.84it/s]
weights: [w1A(11.17), w2A(11.17), wA(6.93)]
passed finite difference test with acc = 1e-08
Trained weights:
Weight 'w1A': 11.172786
Weight 'w2A': 11.173843
Weight 'wA': 6.933418
Testing on OR test-data
test((0.1, 0.1, 0)) returned: 0.000023640720038386 => 0 [correct]
test((0.1, 0.9, 1)) returned: 0.9857913164734841 => 1 [correct]
test((0.9, 0.1, 1)) returned: 0.9057075715094610 => 1 [correct]
test((0.9, 0.9, 1)) returned: 0.999908107780620 => 1 [correct]
Accuracy: 1.000000
-----
Training on AND data
train: 100% | 10000/10000 [00:03:00:00, 3318.70it/s]
weights: [w1A(10.50), w2A(10.50), wA(14.37)]
passed finite difference test with acc = 1e-08
Trained weights:
Weight 'w1A': 10.498804
Weight 'w2A': 10.490561
Weight 'wA': 14.366979
Testing on AND test-data
test((0.1, 0.1, 0)) returned: 4.784254617057318e-06 => 0 [correct]
test((0.1, 0.9, 0)) returned: 0.025484400369173127 => 0 [correct]
test((0.9, 0.1, 0)) returned: 0.02548901063720159 => 0 [correct]
test((0.9, 0.9, 1)) returned: 0.9053604979730813 => 1 [correct]
Accuracy: 1.000000
```

## 6. Two Layer Neural Network:

Using API provided by `neural_net.py` the requested network was implemented.

OR :

```
-----
Training on OR data
train: 100% | 10000/10000 [00:10:00:00, 992.25it/s]
weights: [w1A(-4.17), w2A(-4.15), wA(-2.00), w1B(-5.25), w2B(-5.25), wB(-2.84), w1C(-5.94), w2C(-8.29), wC(-6.10)]
Trained weights:
Weight 'w1A': -4.166633
Weight 'w2A': -4.149468
Weight 'wA': -2.081165
Weight 'w1B': -5.234393
Weight 'w2B': -5.249836
Weight 'wB': -2.835798
Weight 'w1C': -5.038956
Weight 'w2C': -8.290385
Weight 'wC': -6.155050
Testing on OR test-data
test((0.1, 0.1, 0)) returned: 0.000025079020622775 => 0 [correct]
test((0.1, 0.9, 1)) returned: 0.9919010393840599 => 1 [correct]
test((0.9, 0.1, 1)) returned: 0.9919040019425451 => 1 [correct]
test((0.9, 0.9, 1)) returned: 0.9978001861100181 => 1 [correct]
Accuracy: 1.000000
```

AND:

```
.....
Training on AND data
train: 100% | 10000/10000 [00:09:00:00, 1092.36it/s]
weights: [w1A(-3.57), w2A(-5.61), wA(-4.72), w1B(4.02), w2B(3.98), wB(5.31), wAC(-7.77), wBC(8.66), wC(0.96)]
Trained weights:
Weight 'w1A': -3.567641
Weight 'w2A': -5.600579
Weight 'wA': -4.719068
Weight 'w1B': 4.023667
Weight 'w2B': 3.981724
Weight 'wB': 5.306268
Weight 'wAC': -7.765878
Weight 'wBC': 8.655548
Weight 'wC': 0.962835
Testing on AND test-data
test((0.1, 0.1, 0)) returned: 0.89828451642688453718 => 0 [correct]
test((0.1, 0.9, 0)) returned: 0.896783813285854222 => 0 [correct]
test((0.9, 0.1, 0)) returned: 0.896788528452711841 => 0 [correct]
test((0.9, 0.9, 1)) returned: 0.995518965440981 => 1 [correct]
Accuracy: 1.000000
```

EQUAL:

```
.....
Training on EQUAL data
train: 100% | 10000/10000 [00:05:08:08, 1728.78it/s]
weights: [w1A(-6.68), w2A(-6.85), wA(-2.78), w1B(-5.81), w2B(-5.85), wB(-7.58), wAC(18.63), wBC(-10.51), wC(-5.81)]
Trained weights:
Weight 'w1A': -6.682529
Weight 'w2A': -6.845288
Weight 'wA': -2.782586
Weight 'w1B': -5.812516
Weight 'w2B': -5.856722
Weight 'wB': -7.581487
Weight 'wAC': 18.631200
Weight 'wBC': -10.505650
Weight 'wC': -5.800065
Testing on EQUAL test-data
test((0.1, 0.1, 1)) returned: 0.9572829757678457 => 1 [correct]
test((0.1, 0.9, 0)) returned: 0.811226265758113587 => 0 [correct]
test((0.9, 0.1, 0)) returned: 0.811358471362683439 => 0 [correct]
test((0.9, 0.9, 1)) returned: 0.9683879694444555 => 1 [correct]
Accuracy: 1.000000
```

HORIZONTAL BAND:

```
.....
Training on horizontal-bands data
train: 100% | 10000/10000 [00:24:00:00, 405.12it/s]
weights: [w1A(0.87), w2A(-0.67), wA(-11.30), w1B(0.22), w2B(-0.74), wB(-3.21), wAC(12.12), wBC(-11.04), wC(5.90)]
Trained weights:
Weight 'w1A': 0.874565
Weight 'w2A': -0.668671
Weight 'wA': -11.384945
Weight 'w1B': 0.210894
Weight 'w2B': -0.737621
Weight 'wB': -3.212918
Weight 'wAC': 12.123274
Weight 'wBC': -11.838996
Weight 'wC': 5.903178
Testing on horizontal-bands test-data
test((1, 1.5, 1)) returned: 0.9978851982897587 => 1 [correct]
test((2, 1.5, 1)) returned: 0.9978995487086657 => 1 [correct]
test((3, 1.5, 1)) returned: 0.9977181878957527 => 1 [correct]
test((0, 1.5, 1)) returned: 0.9976673543307334 => 1 [correct]
test((4, 0, 0)) returned: 0.884398324808839863 => 0 [correct]
test((4, 4, 0)) returned: 0.8827537375878151226 => 0 [correct]
test((-1, 0, 0)) returned: 0.886341633795682344 => 0 [correct]
test((-1, 4, 0)) returned: 0.8827442397577797715 => 0 [correct]
Accuracy: 1.000000
```

## VERTICAL BAND:

```
.....
Training on vertical-bands data
train:: 100% | 10000/10000 [00:24<00:00, 411.54it/s]
weights: [w1A(-4.63), w2A(0.07), wA(-11.29), w1B(-6.68), w2B(0.19), wB(-3.28), wAC(12.17), wBC(-11.88), wC(5.92)]
Trained weights:
Weight 'w1A': -4.631307
Weight 'w2A': 0.073578
Weight 'wA': -11.298081
Weight 'w1B': -6.682839
Weight 'w2B': 0.194834
Weight 'wB': -3.196799
Weight 'wAC': 12.171234
Weight 'wBC': -11.884988
Weight 'wC': 5.921322
Testing on vertical-bands test-data
test((0, 1, 0)) returned: 0.005238871070315181 => 0 [correct]
test((0, 2, 0)) returned: 0.004896216045513198 => 0 [correct]
test((0, 1.5, 0)) returned: 0.005853170168565381 => 0 [correct]
test((1.5, 2, 1)) returned: 0.9977478890786847 => 1 [correct]
test((1.5, 5, 1)) returned: 0.9977603400740385 => 1 [correct]
test((1.5, 1, 1)) returned: 0.9977333610340835 => 1 [correct]
test((3, 1, 0)) returned: 0.00654137323795754 => 0 [correct]
test((3, 1.5, 0)) returned: 0.006746413743868805 => 0 [correct]
test((3, 2, 0)) returned: 0.006664631731464771 => 0 [correct]
test((1, 1.5, 1)) returned: 0.9967653131579071 => 1 [correct]
test((1, -1.5, 1)) returned: 0.9973872983300298 => 1 [correct]
test((2, 1.5, 1)) returned: 0.9938786877614689 => 1 [correct]
test((2, -1.5, 1)) returned: 0.9989087884424382 => 1 [correct]
test((4, 0, 0)) returned: 0.002697970682214445 => 0 [correct]
test((4, 4, 0)) returned: 0.0027868538424612276 => 0 [correct]
test((-1, 0, 0)) returned: 0.00356803740131587 => 0 [correct]
test((-1, 4, 0)) returned: 0.0035587754638094517 => 0 [correct]
Accuracy: 1.000000
```

## DIAGONAL BAND:

```
.....
Training on diagonal-band data
train:: 100% | 10000/10000 [00:25<00:00, 392.95it/s]
weights: [w1A( 4.37), w2A(4.76), wA( 4.05), w1B(4.76), w2B( 4.43), wB( 4.14), wAC(18.04), wBC(18.03), wC(15.12)]
Trained weights:
Weight 'w1A': 4.367353
Weight 'w2A': 4.761861
Weight 'wA': -4.049311
Weight 'w1B': 4.750038
Weight 'w2B': -4.434639
Weight 'wB': -4.141418
Weight 'wAC': 18.036732
Weight 'wBC': 18.033377
Weight 'wC': 15.122426
Testing on diagonal-band test-data
test(( 1, -1, 1)) returned: 0.9889386317928835 => 1 [correct]
test((5, 5, 1)) returned: 0.9924861819778131 => 1 [correct]
test((-2, -2, 1)) returned: 0.9888919510788497 => 1 [correct]
test((0, 0, 1)) returned: 0.992619064389892 => 1 [correct]
test((3.5, 3.5, 1)) returned: 0.9921754795060245 => 1 [correct]
test((1.5, 1.5, 1)) returned: 0.9915053013937995 => 1 [correct]
test((0, 0, 0)) returned: 0.0061262124552528184 => 0 [correct]
test((8, 4, 0)) returned: 0.006146656123367745 => 0 [correct]
Accuracy: 1.000000
```

## INVERSE DIAGONAL BAND:

```
Training on inverse diagonal band data
train: 100% | 10000/10000 [00:24<00:00, 412.30it/s]
weights: [w1A(-4.32), w2A(4.01), w1B(-5.95), w1B(4.02), w2B(-4.35), w1C(-5.95), w1C(10.00), w2C(10.00), wC(5.10)]
Trained weights:
Weight 'w1A': -4.323526
Weight 'w2A': 4.009196
Weight 'w1B': 3.947959
Weight 'w1B': 4.016341
Weight 'w2B': -4.327655
Weight 'w1C': 3.961804
Weight 'w1C': 10.796815
Weight 'w2C': 10.796226
Weight 'wC': 5.176297
Testing on inverse-diagonal-band test data
test((-1, -1, 0)) returned: 0.009700016325943552 => 0 [correct]
test((5, 5, 0)) returned: 0.00611940638429841 => 0 [correct]
test((-2, -2, 0)) returned: 0.011770625765937533 => 0 [correct]
test((6, 6, 0)) returned: 0.005981083944559856 => 0 [correct]
test((3.5, 3.5, 0)) returned: 0.0064401057268764435 => 0 [correct]
test((1.5, 1.5, 0)) returned: 0.007242086015958169 => 0 [correct]
test((4, 0, 1)) returned: 0.9963879772177346 => 1 [correct]
test((0, 4, 1)) returned: 0.9963879772177346 => 1 [correct]
Accuracy: 1.000000
```


## NOT EQUAL:

```
Training on NOT EQUAL data
train: 100% | 10000/10000 [00:05<00:00, 1743.43it/s]
weights: [w1A(-6.58), w2A(6.50), w1B(3.50), w1B(-5.97), w2B(6.23), w1C(-2.97), w1C(10.22), w2C(-10.01), wC(-4.70)]
Trained weights:
Weight 'w1A': -6.582623
Weight 'w2A': 6.587322
Weight 'w1B': 3.588473
Weight 'w1B': -5.969659
Weight 'w2B': 6.238543
Weight 'w1C': 2.973468
Weight 'w1C': 10.221127
Weight 'w2C': 10.006440
Weight 'wC': 4.788817
Testing on NOT EQUAL test data
test((0.1, 0.1, 0)) returned: 0.011342505438581542 => 0 [correct]
test((0.1, 0.9, 1)) returned: 0.9078214787819529 => 1 [correct]
test((0.9, 0.1, 1)) returned: 0.9058289458612442 => 1 [correct]
test((0.9, 0.9, 0)) returned: 0.010428404981705212 => 0 [correct]
Accuracy: 1.000000
```




## 7. Plot Decision Boundary:


To plot the corresponding decision boundary `plt.scatter` was used.  
NOT EQ boundaries:



OR boundaries:



AND boundaries:



## 8. Overfitting and Regularization:

In short, Regularization in machine learning is the process of regularizing the parameters that constrain, regularizes, or shrinks the coefficient estimates towards zero. In other words, this technique discourages learning a more complex or flexible model, avoiding the risk of Overfitting.

Iteration = 100

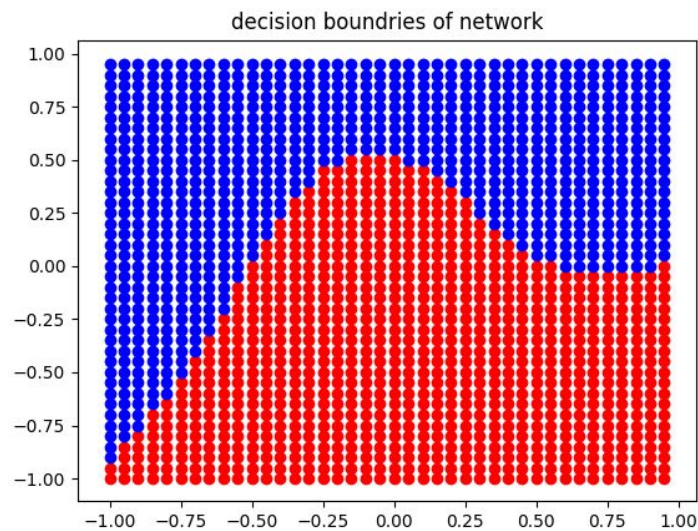
```
* code_sajjad python3 neural_net_tester.py two_moons
```

```
.....  
Training on two moons data
```

```
train:: 100%  
decision surface:: 1600it [00:07, 219.07it/s]
```

```
Accuracy: 0.990000
```

```
| 100/100 [00:25<00:00, 3.86it/s]
```



Iteration = 500

```
* code_sajjad python3 neural_net_tester.py two_moons
```

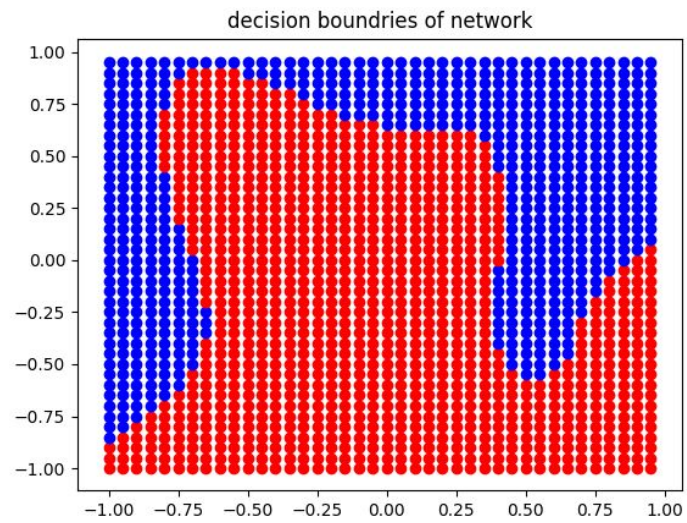
```
.....  
Training on two moons data
```

```
train:: 100%  
decision surface:: 1600it [00:08, 178.69it/s]
```

```
Accuracy: 0.990000
```


```
| 500/500 [02:13<00:00, 3.74it/s]
```





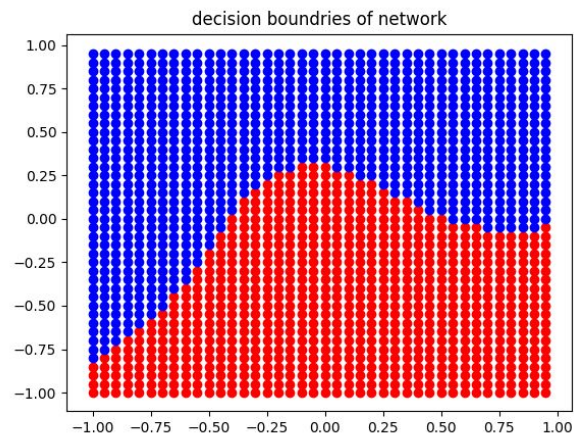
Iteration 1000:

```
-----  
Training on two-moons data  
train: 100% | 1000/1000 [04:36<00:00, 3.50it/s]  
decision surface:: 1600it [00:00, 182.38it/s]  
Accuracy: 0.862000
```




Iteration = 100 with regularization:

```
* code_wajjad python3 neural_net_tester.py two_moons
-----
Training on two-moons data
train: 100% | 100/100 [00:27<00:00, 3.59it/s]
decision surface: 1600it [00:07, 227.40it/s]
Accuracy: 0.980000
```



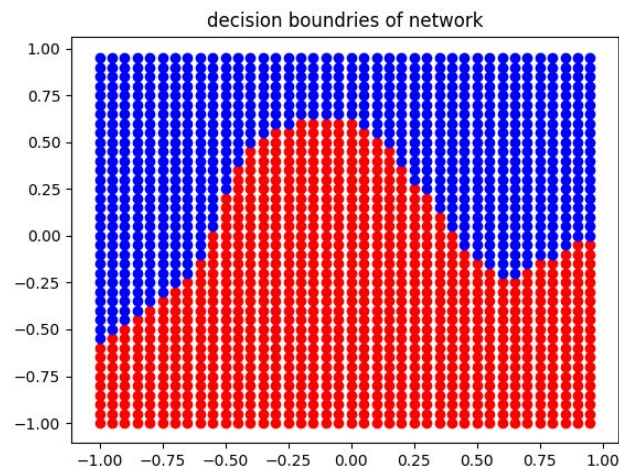
Iteration = 500 with regularization:

```
* code_wajjad python3 neural_net_tester.py two_moons
-----
Training on two-moons data
train: 100% | 500/500 [02:28<00:00, 3.57it/s]
decision surface: 1600it [00:07, 227.70it/s]
Accuracy: 0.980000
```



Iteration = 1000 with regularization:

```
+ code_sajjad python3 neural_net_tester.py two moons
-----
Training on two-moons data
train:: 100% | 1000/1000 | 04:57:00:00, 3.36it/s |
decision surface:: 1000it | 00:10, 157.11it/s |
Accuracy: 0.980808
```



In the first part as the number of iterations increased , model was overfitted on training data. It happens mainly for the extreme flexibility of neural networks. One of the signs of overfitting is complex decision boundaries which occurred in this part.

When regularization term is added , it makes parameters to constantly drift to 0 thus prevents overfitting by putting additional yet desirable constraints. As a result decisions boundaries would appear more smooth.

The problem of overfitting is obviously solved!